SOLVCREST AI

Market Landscape Review

From a market summary to an interview-ready validation plan.

PART 1 **Mentor review**

PART 2 **Revised landscape**

Prepared for Tom Gousby

Founder Institute mentor review | October 18, 2026

Items marked Verify or ? are hypotheses to confirm through research or customer interviews.

EXECUTIVE SUMMARY

Strong founder insight. Unproven demand.

This is a clean one-pager, but it reads as a market summary, not validation. It's honest about its limits, and that's the best thing in it. Part 1 shows where the landscape falls short. Part 2 replaces it with an interview-ready version: positioning, beachhead, bottom-up sizing, a competitor matrix, and a hypothesis tracker to break.

0
customer interviews completed so far

7
hypotheses to prove or kill

30
discovery calls targeted in 60 days

~670
customers needed for $100M ARR at a $150k contract

Mentor scorecard (out of 10)

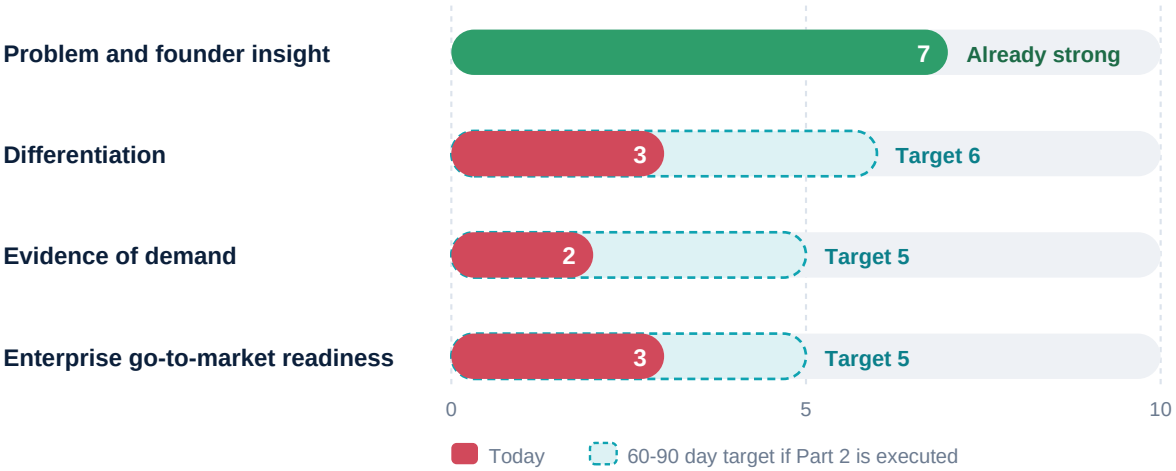

Scores are mentor judgment, not measurements. Venture potential if the idea is narrowed and validated: **6-7 out of 10**.

How this report is organized

PART 1

Mentor review

- What's good
- Six gaps
- Fixes at a glance

PART 2

Revised landscape

1. Positioning
2. Beachhead
3. Bottom-up market
4. Competitor matrix
5. Hypothesis tracker
6. Interview kit
7. Evidence ladder
8. Before the first call

PART 1

Mentor review

What the original one-pager gets right, and where it falls short

What's good

1 Honest labeling

"Not Solvcrest's directly serviceable market," "survey-based, not predicted ROI," and "no primary customer interviews completed yet" are candid caveats. Keep that tone.

2 The right validation questions

How often incidents repeat, who owns the budget, and whether a buyer would authorize a paid pilot is your whole thesis.

3 It names the incumbent threat

The one-pager faces the competition instead of hiding it.

Where it falls short

1 The market number does nothing for you

A $36B AIOps platform market is top-down and irrelevant. Nobody at Founder Institute or on an investment committee will credit a slice of that. They want bottom-up: how many target companies, how many teams, at what contract value.

Fixed in Part 2, section 3

2 The $1.8M/hour stat measures the wrong thing

That is the cost of a high-impact outage. Your product prevents *repeat* incidents, so you need the cost of recurrence: the share of incidents that repeat and what a repeat costs. Verify the exact figure, wording, and report year; neither source has been checked.

Fixed in Part 2, sections 5 and 8

3 The competitor list is thin, with no verdict

Four AI SRE startups, two incident-management tools, and two observability incumbents, then "formidable." It never says why a buyer picks you, and it misses the competitors closest to your pipeline.

Fixed in Part 2, section 4

4 The "beachhead" isn't a beachhead

"Regulated organizations with Java, Python, GraphQL, and API-heavy applications" describes a tech stack, not a customer. A beachhead names a buyer, a company type, a trigger, and a reason they buy first.

Fixed in Part 2, section 2

5 There is no path to venture scale

At a $150k average contract, a $100M ARR business needs roughly 670 customers. Show where they come from: other regulated sectors, adjacent workflows, or expansion within each account.

Fixed in Part 2, section 3

6 It proves the problem, not your right to win

It validates that the problem exists. Nothing in it is evidence from a customer.

Fixed in Part 2, sections 5-7

The competitors you are missing

These sit closest to your actual pipeline, the ones that go from error to code to a proposed fix.

Error-monitoring AI fix agents

Tools with built-in AI fix agents, such as Sentry's Seer

Observability coding agents

Observability vendors' coding agents, such as Datadog's

General coding agents

Claude Code, Codex, or Copilot's agent pointed at a stack trace

MENTOR NOTE

Your pipeline, trace to file and blame to fix, is territory those tools already cover. Verify the current state of each. **Your edge has to be the history and recurrence layer.**

A better beachhead sentence

Production-support and problem-management teams at mid-to-large U.S. financial firms with in-house Java/Python services and a ServiceNow incident history.

Fixes at a glance

BEFORE	AFTER
Top-down market: $36B AIOps platform market	**Bottom-up SAM: firms x teams x contract value**
Outage-cost stat: $1.8M per hour, high-impact outages	**Recurrence cost: repeat share and cost per repeat, from interviews**
Beachhead as tech stack: Java, Python, GraphQL, APIs	**Beachhead as customer: buyer, trigger, data prerequisite**
Landscape one-pager with no customer evidence	**Hypothesis tracker with kill thresholds and an evidence column**
Competition called "formidable", no verdict	**Scored matrix, a defensible claim, and a kill-test**

THE HEADLINE TO EARN

In financial-services production support, **[X]%** of incidents are repeats, costing **[$Y]** per team a year. No tool today closes the loop from diagnosis to verified permanent fix. *Do not publish this sentence until your interviews support it.*

PART 2

Revised landscape, interview-ready

Replaces the original one-pager

Everything marked **Verify** or **?** is a hypothesis to fill from research or interviews. Nothing in the tracker counts as evidence until a source or a named customer is attached.

1

PART 2

Positioning

POSITIONING

Solvcrest AI is the incident-memory and permanent-fix engine for production-support teams.

It traces each failure to the exact file, function, and commit, recognizes whether it is new, similar, recurring, or a regression, and drafts a human-approved fix, so repeat incidents get eliminated instead of re-resolved.

TODAY: WORKING PROTOTYPE (FROM THE REPO)

Log
input
Python / Java traces, GraphQL errors

→

Parser
parser.py
Traces become a structured signal

→

Locator
locator.py
Repo file, snippet, git blame, recent commits

→

History
history.py
Error signature plus similarity search in SQLite

→

Diagnosis
agent.py / rules.py
LLM analysis, or rules fallback with no key

→

Report
report.py
Human-readable and JSON output

History verdict

NEW
No prior match in history

SIMILAR
Reuse fixes that worked on near matches

RECURRING
Same signature, only workarounds recorded: fix the root cause

REGRESSION
Previously fixed permanently: look at what changed

Close the loop after every incident: record the root cause and permanent fix so the next match learns.

NOT BUILT YET (REPO NEXT STEPS)

- Webhook ingestion: Datadog, Loki, CloudWatch, Sentry, PagerDuty
- Embeddings for history, plus deploy and metric correlation
- Open a PR with the patch; human approval, never auto-merge
- Eval harness: replay past incidents, score fix-location accuracy
- Guardrails: read-only repo access, secret scrubbing

PART 2

2 Beachhead

Company type	Mid-to-large U.S. financial firms (banks, asset managers, insurers) with in-house Java/Python services
Team	Production support, application support, SRE, and problem-management teams
Champion	SRE lead or production-support manager who owns recurring incidents
Economic buyer	Head of production services or engineering ([title to confirm])
Blockers	Infosec, vendor risk, model risk, change management
Trigger	A repeat Sev-1/Sev-2, an audit finding, or a postmortem action item that never closed
Data prerequisite	Incident history in ServiceNow or Jira, plus read access to the service repos
Why first	Regulated buyers feel recurrence pain hardest and need audit trails that general tools skip

How a deal moves through the buying committee

TRIGGER
Repeat Sev-1 / Sev-2, an audit finding, or a postmortem action that never closed

→

CHAMPION
SRE lead or production-support manager who owns recurring incidents

→

ECONOMIC BUYER
Head of production services or engineering

→

REVIEW GATES
- Infosec
- Vendor risk
- Model risk
- Change mgmt

→

PILOT
Read-only, customer-hosted, success criteria agreed up front

PART 2

3 Bottom-up market table

Fill the **?** boxes from research and interviews. The scenario strip is illustrative arithmetic only, not a market claim.

TARGET FIRMS		TEAMS PER FIRM		ACV PER TEAM		SAM
?	x	?	x	?	=	?
Your research		Interviews		Pilot pricing test		Firms x teams x ACV

ILLUSTRATIVE ONLY
Arithmetic, not research

300 firms x 3 teams x $100k = $90M

CUSTOMERS NEEDED FOR $100M ARR

Customers = $100M / annual contract value (arithmetic)

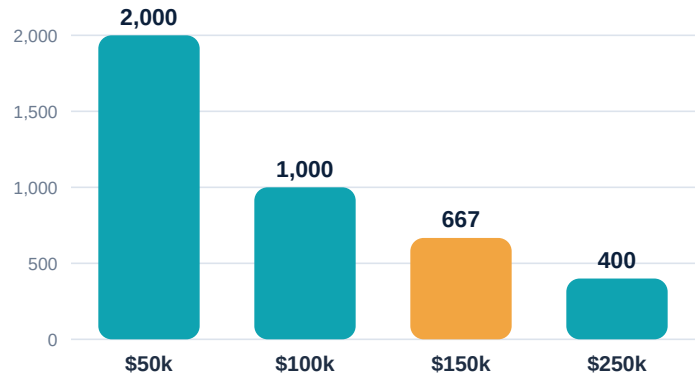

average annual contract value (amber: the $150k case from Part 1)

THE BEACHHEAD ALONE LIKELY FALLS SHORT

Name the expansion path now:

1. **Other regulated sectors**
   Insurance, healthcare, public sector
2. **More teams per firm**
   Land one team, expand across the firm
3. **Adjacent workflows**
   For example, change-risk review

PART 2

4 Competitor matrix

Score each cell with a source link or a demo you watched. **Likely** means an unverified guess to check. Be honest in the last column about where you lose today.

Player	Cross-incident memory	Recurrence tracking	Code-level fix location	Verified-fix loop	Customer-hosted	Where they likely win
Solvcrest (today, per repo)	**Yes (local SQLite)**	**Yes (new / similar / recurring / regression)**	**Yes (file:line, git blame; Python and Java only)**	**Not yet (manual resolve, no PR generation)**	**Yes (local-first)**	Not yet: no integrations, no eval accuracy number
AI SRE specialists (Resolve AI, Cleric, Ciroos, Traversal)	Verify	Verify	Verify	Verify	Verify	Likely breadth, funding, integrations
Observability incumbents (Datadog, Dynatrace)	Verify	Verify	Verify	Verify	Verify	Likely existing telemetry access and bundling
Incident management (Rootly, PagerDuty)	Verify	Verify	Verify	Verify	Verify	Likely workflow ownership and installed base
Error-monitoring AI fix agents (for example Sentry Seer)	Verify	Verify	**Likely strong**	Verify	Verify	Likely error-to-code-to-fix on a single incident
Generic coding agent (Claude Code, Codex, Copilot agent) on a pasted log	**Likely none**	**Likely none**	**Likely strong**	**Likely none**	Verify	Likely speed and price for one-off diagnosis

Confirmed in the repo · Gap today · Likely, unverified · Verify with a source

The wedge, section 4 continued

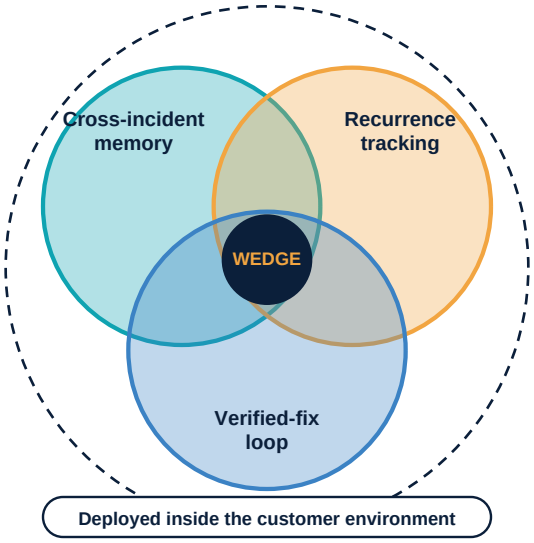

YOUR DEFENSIBLE CLAIM, IF THE RESEARCH HOLDS

Cross-incident memory plus **recurrence tracking** plus a **verified-fix loop**, deployable inside the customer's environment. Single-incident diagnosis is not a moat.

KILL-TEST

Run the same 20 incidents through Solvcrest and through a generic coding agent with no history. If the agent matches you on repeat detection, rethink the wedge.

PART 2

5 Hypothesis tracker

Seven bets, each with a test that can kill it. Attach a call date and a named contact to any evidence you log.

#	Hypothesis	Disproof test	Kill threshold	Evidence	Status
H1	**A meaningful share of incidents at target firms are repeats**	Recurrence Audit on sanitized ticket data from 3+ firms	Under [X]% repeats	None	Untested
H2	**Repeats cost more than [$Y] per quarter per team**	Interviews, quantified	Cannot quantify in 7 of 10 calls	None	Untested
H3	**Incumbent tools do not close the loop to a verified permanent fix**	Side-by-side on the same incidents	Incumbent matches on 80%+	None	Untested
H4	**A named owner has budget for recurrence elimination**	30 discovery calls	No budget owner in 20 of 30	None	Untested
H5	**Buyers will pay for a pilot**	Paid-pilot or LOI ask	0 of 10 qualified accounts agree	None	Untested
H6	**Security and model-risk review can be cleared with a read-only, customer-hosted design**	Security packet review with 2+ prospects	Both reject the architecture	None	Untested
H7	**The in-house-code slice of incidents is large enough**	Ask each interviewee what share of incidents trace to code they own	Under [X]%	None	Untested

STATUS VALUES: Untested | Supported | Weakened | Killed

PART 2

6 Interview kit

Who to book: 60 targets, 30 calls

Heads of production support, SRE and app-support managers, and problem-management leads at financial firms. Start with your own network and the Founder Institute community.

RULE Do not demo in the first call

RULE Keep it to 25 minutes

Seven questions

1. Walk me through your last three repeat incidents. What happened each time?
2. What did each one cost in hours, dollars, and customer impact?
3. Who owned the permanent fix, and why didn't it ship the first time?
4. Where does your incident history live today, and who can query it?
5. What do you already use for AI-assisted diagnosis, and where does it fall short?
6. Who would approve a pilot, and what would security and risk need to see?
7. What do you spend on this today, in tools and in people-hours?

Score every call

0 - 2	0 - 2	0 - 2	0 - 2	0 - 2
Pain quantified	Budget owner identified	Repeats confirmed	Data access plausible	Next step agreed

Under 5: stays a contact, not a lead

5 or more: qualified lead

0 1 2 3 4 5 6 7 8 9 10

The Recurrence Audit offer

THEY SEND

A sanitized export of 6-12 months of incident tickets

Recurrence Audit

No security review needed. Tests H1, H2, H4.

YOU RETURN

Top repeat offenders, estimated cost, and a ranked permanent-fix backlog

RED FLAGS

1. "Interesting" with no next step. **2.** No budget owner. **3.** Incidents blamed entirely on vendors or infrastructure you cannot touch. **4.** Refusal to share even sanitized data.

PART 2

7 Evidence ladder

Targets for the next 60 days. Fill the Actual column only with rungs someone actually climbed. A paid pilot is worth more than a hundred friendly calls.

EVIDENCE RUNG	TARGET	ACTUAL
Discovery interview completed	30	0
Pain quantified in dollars or hours	10	0
Sanitized data shared	5	0
Design partner (named contact, weekly call)	3	0
Signed pilot LOI	2	0
Paid pilot ($10-25k)	1	0

Count only rungs someone actually climbed.

A suggested 60-day path

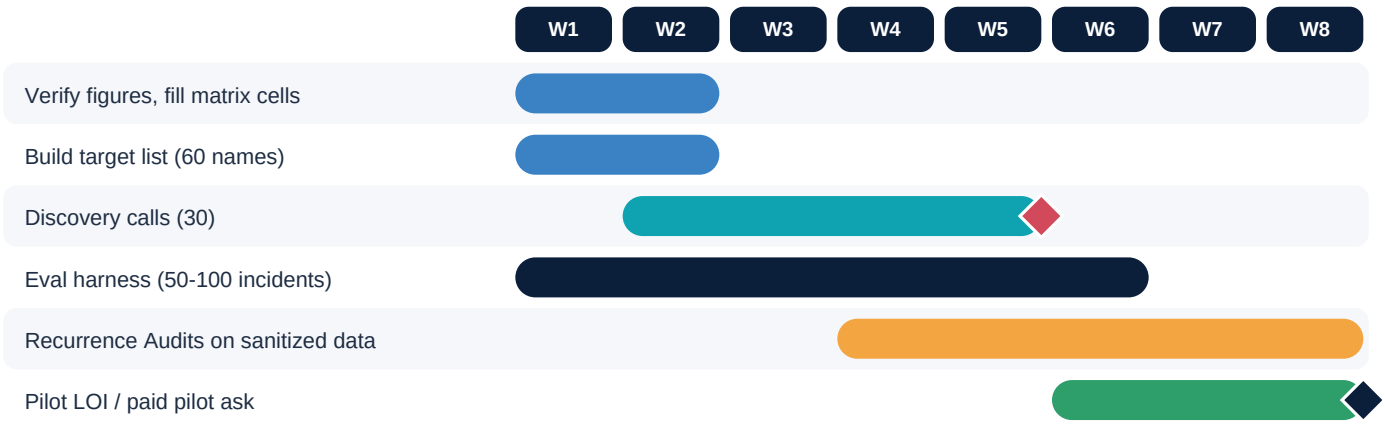

Suggested sequencing for the 60-day window; adjust to your calendar. Diamonds mark the 30-call and first-paid-pilot targets.

PART 2

8 Before the first call

- [ ] Verify the Grand View Research and New Relic figures against the sources, including report year and exact wording, or drop them.
- [ ] Replace the $1.8M-per-hour outage stat with a recurrence cost you derive from interviews.
- [ ] Fill every **Verify** cell in the competitor matrix from product pages or demos, with links.
- [ ] Build the eval harness: replay 50-100 real or public incidents and score fix-location accuracy and repeat detection against a generic coding-agent baseline.
- [ ] Keep data clean: use only sanitized or synthetic data and your own non-employer contacts, and resolve any employer IP or outside-activity questions before accepting a pilot.

SOURCES NAMED IN THE ORIGINAL ONE-PAGER, TO VERIFY

grandviewresearch.com/industry-analysis/aiops-market | newrelic.com/resources/report/observability-forecast

THE RULE OF THIS REPORT

Count only rungs someone actually climbed.

Update the tracker after every call, and attach a call date and a named contact to any evidence.